

Commander

Better Distributed Applications through CQRS, Event Sourcing, and Immutable Logs

Hi, I'm Bobby

I'm on the Technology Fellow's team at  **Capital One**

I dislike accidental complexity

bobby.calderwood@capitalone.com

[@bobbycalderwood](#)

<https://github.com/bobby>


Problem Space

- Provide valuable informational and transactional services via Mobile and Web software
- To ***lots*** of customers, with excellent user experience
- Securely and in compliance with regulations
- With ability to easily enhance, experiment, monitor, maintain, and operate
- By many participants within a large organization

Big Ideas

- Immutability is central to information systems
- Data language of system >> Programming language of components
- Action and perception are not the same, and immutability facilitates their separation
- Businesses services are not databases, they're event stream reactors
- Cross-cutting concerns must be satisfied in the presence of Conway's Law

Problematic Architecture




Immutability is central
to information systems

Event Sourcing Analogy



Image by [Alan Light](#) CC BY-SA 3.0

Data Loss by Design

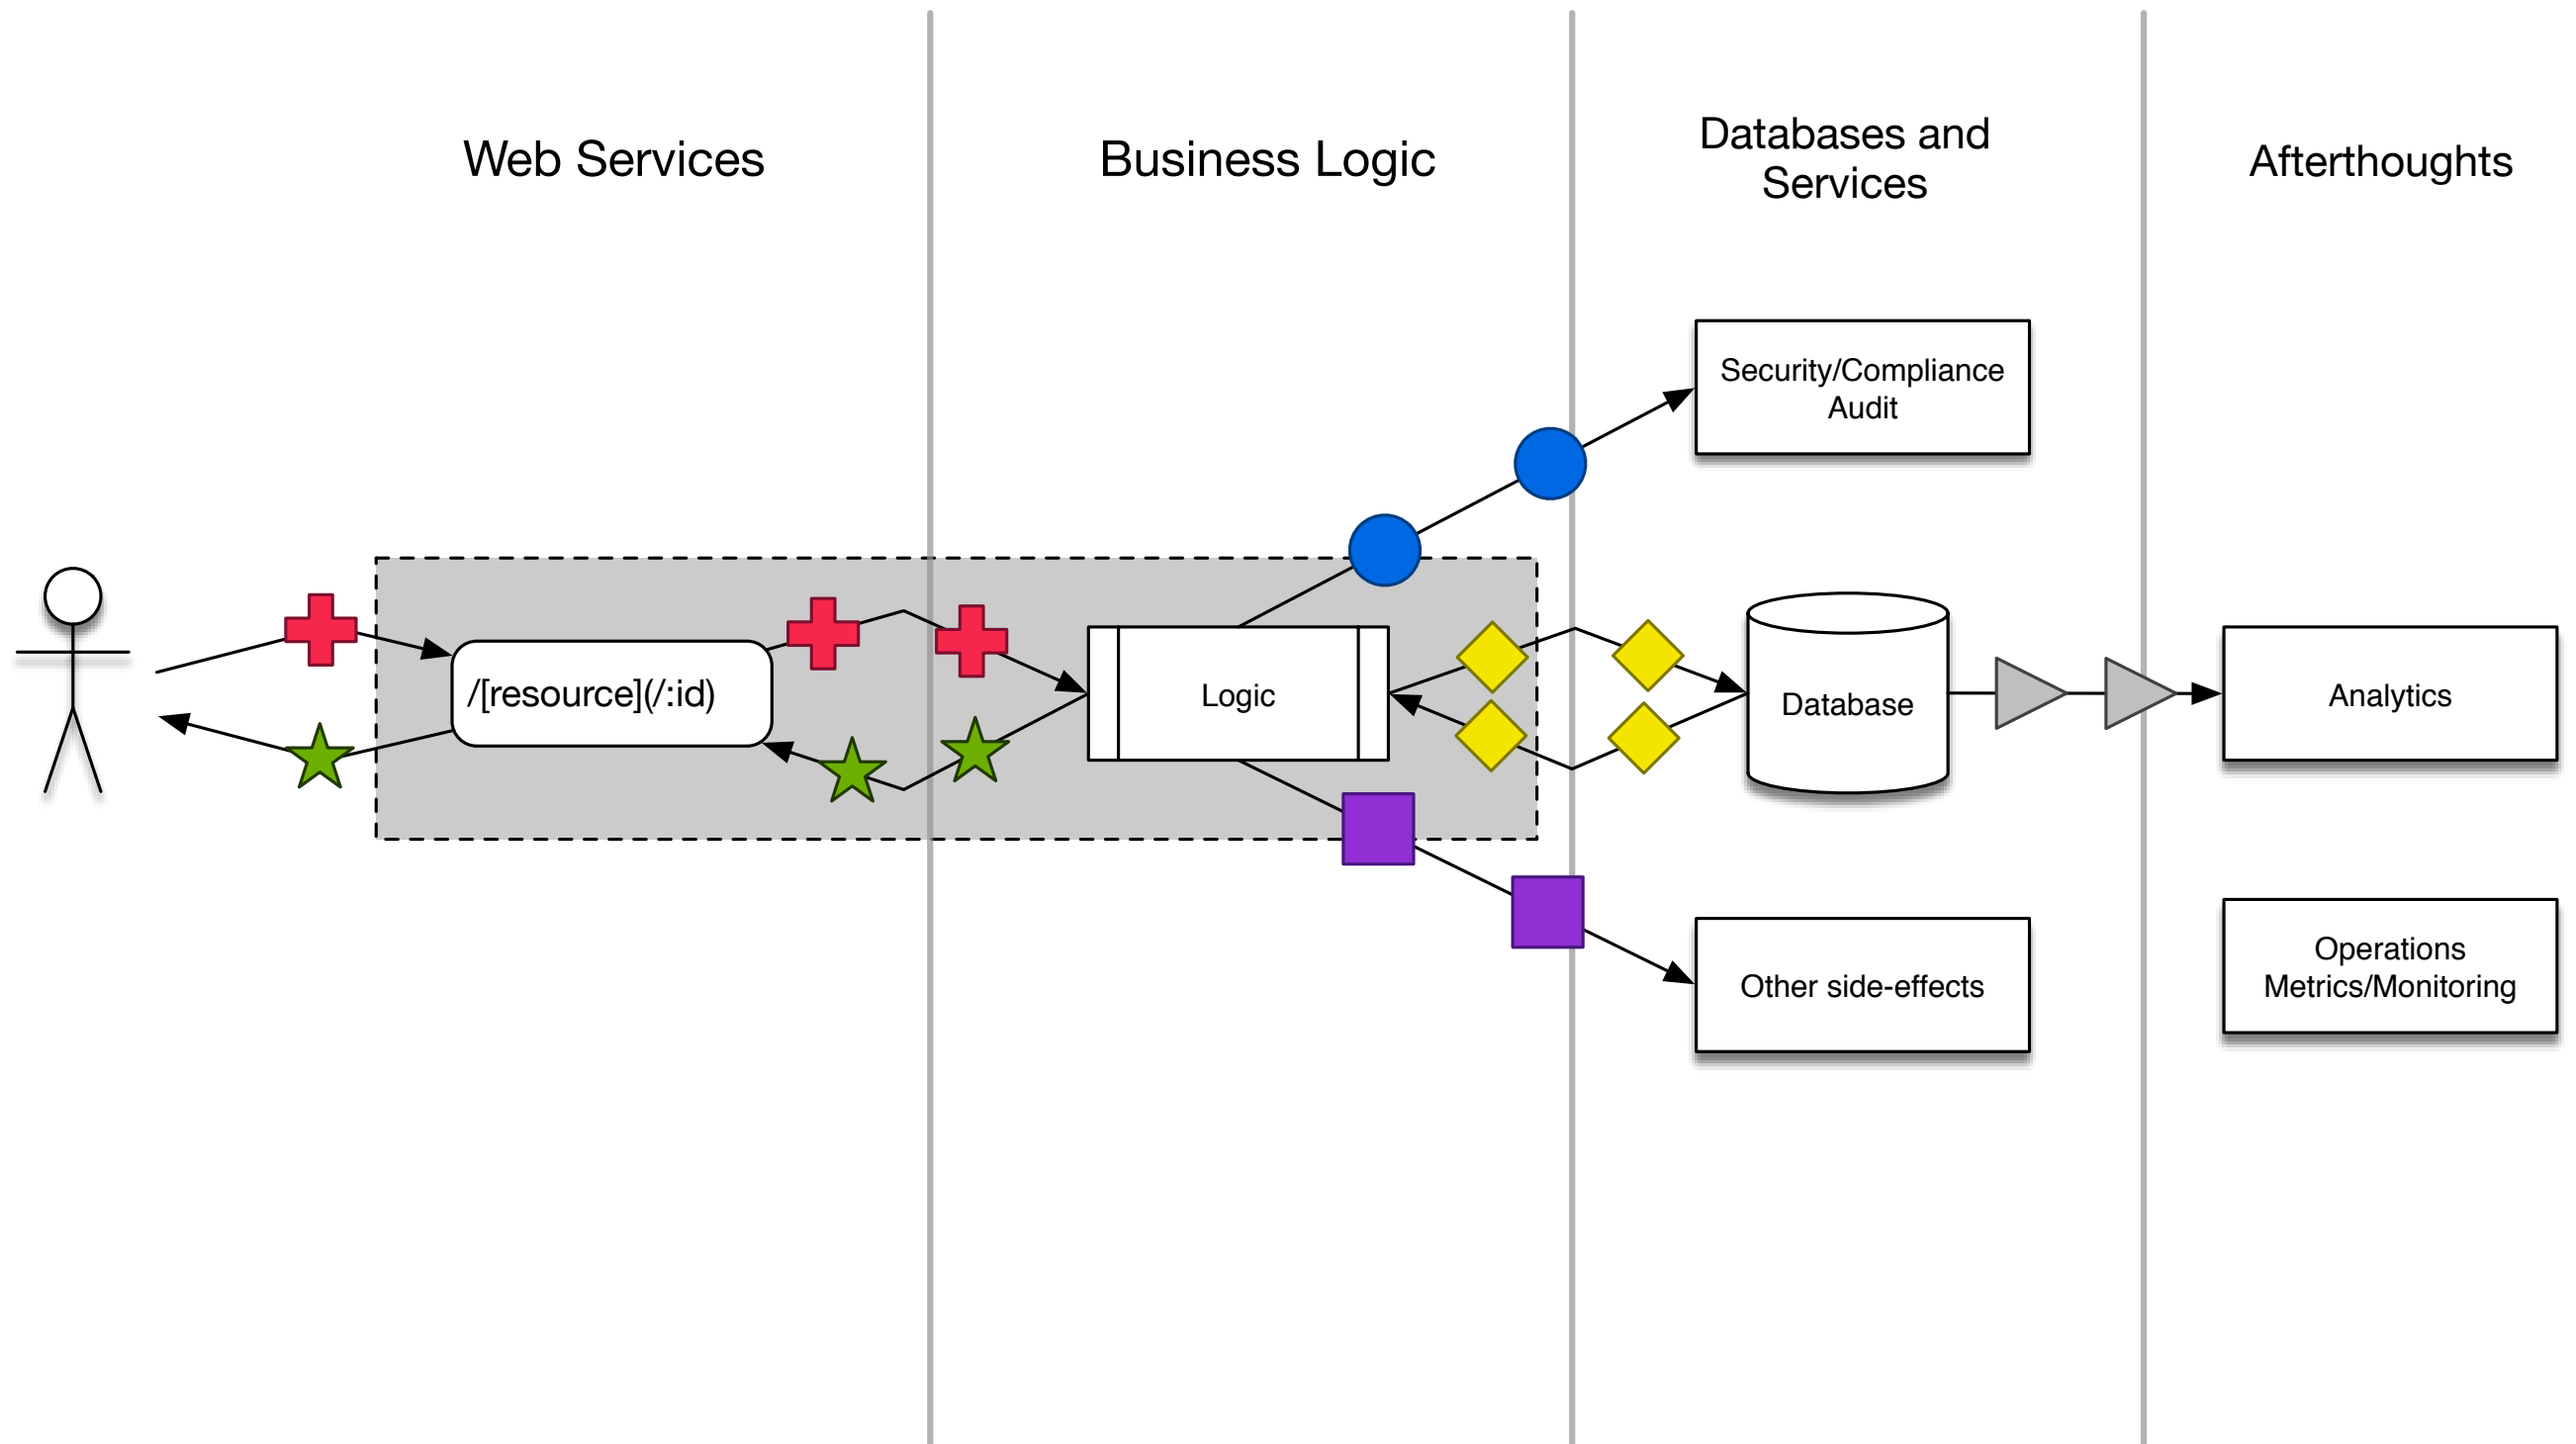


The (data) language of the System

>>

The (runtime) language of each
component


What is system language?



Action \neq Perception

Writes \neq Reads

Writes Tied to Reads

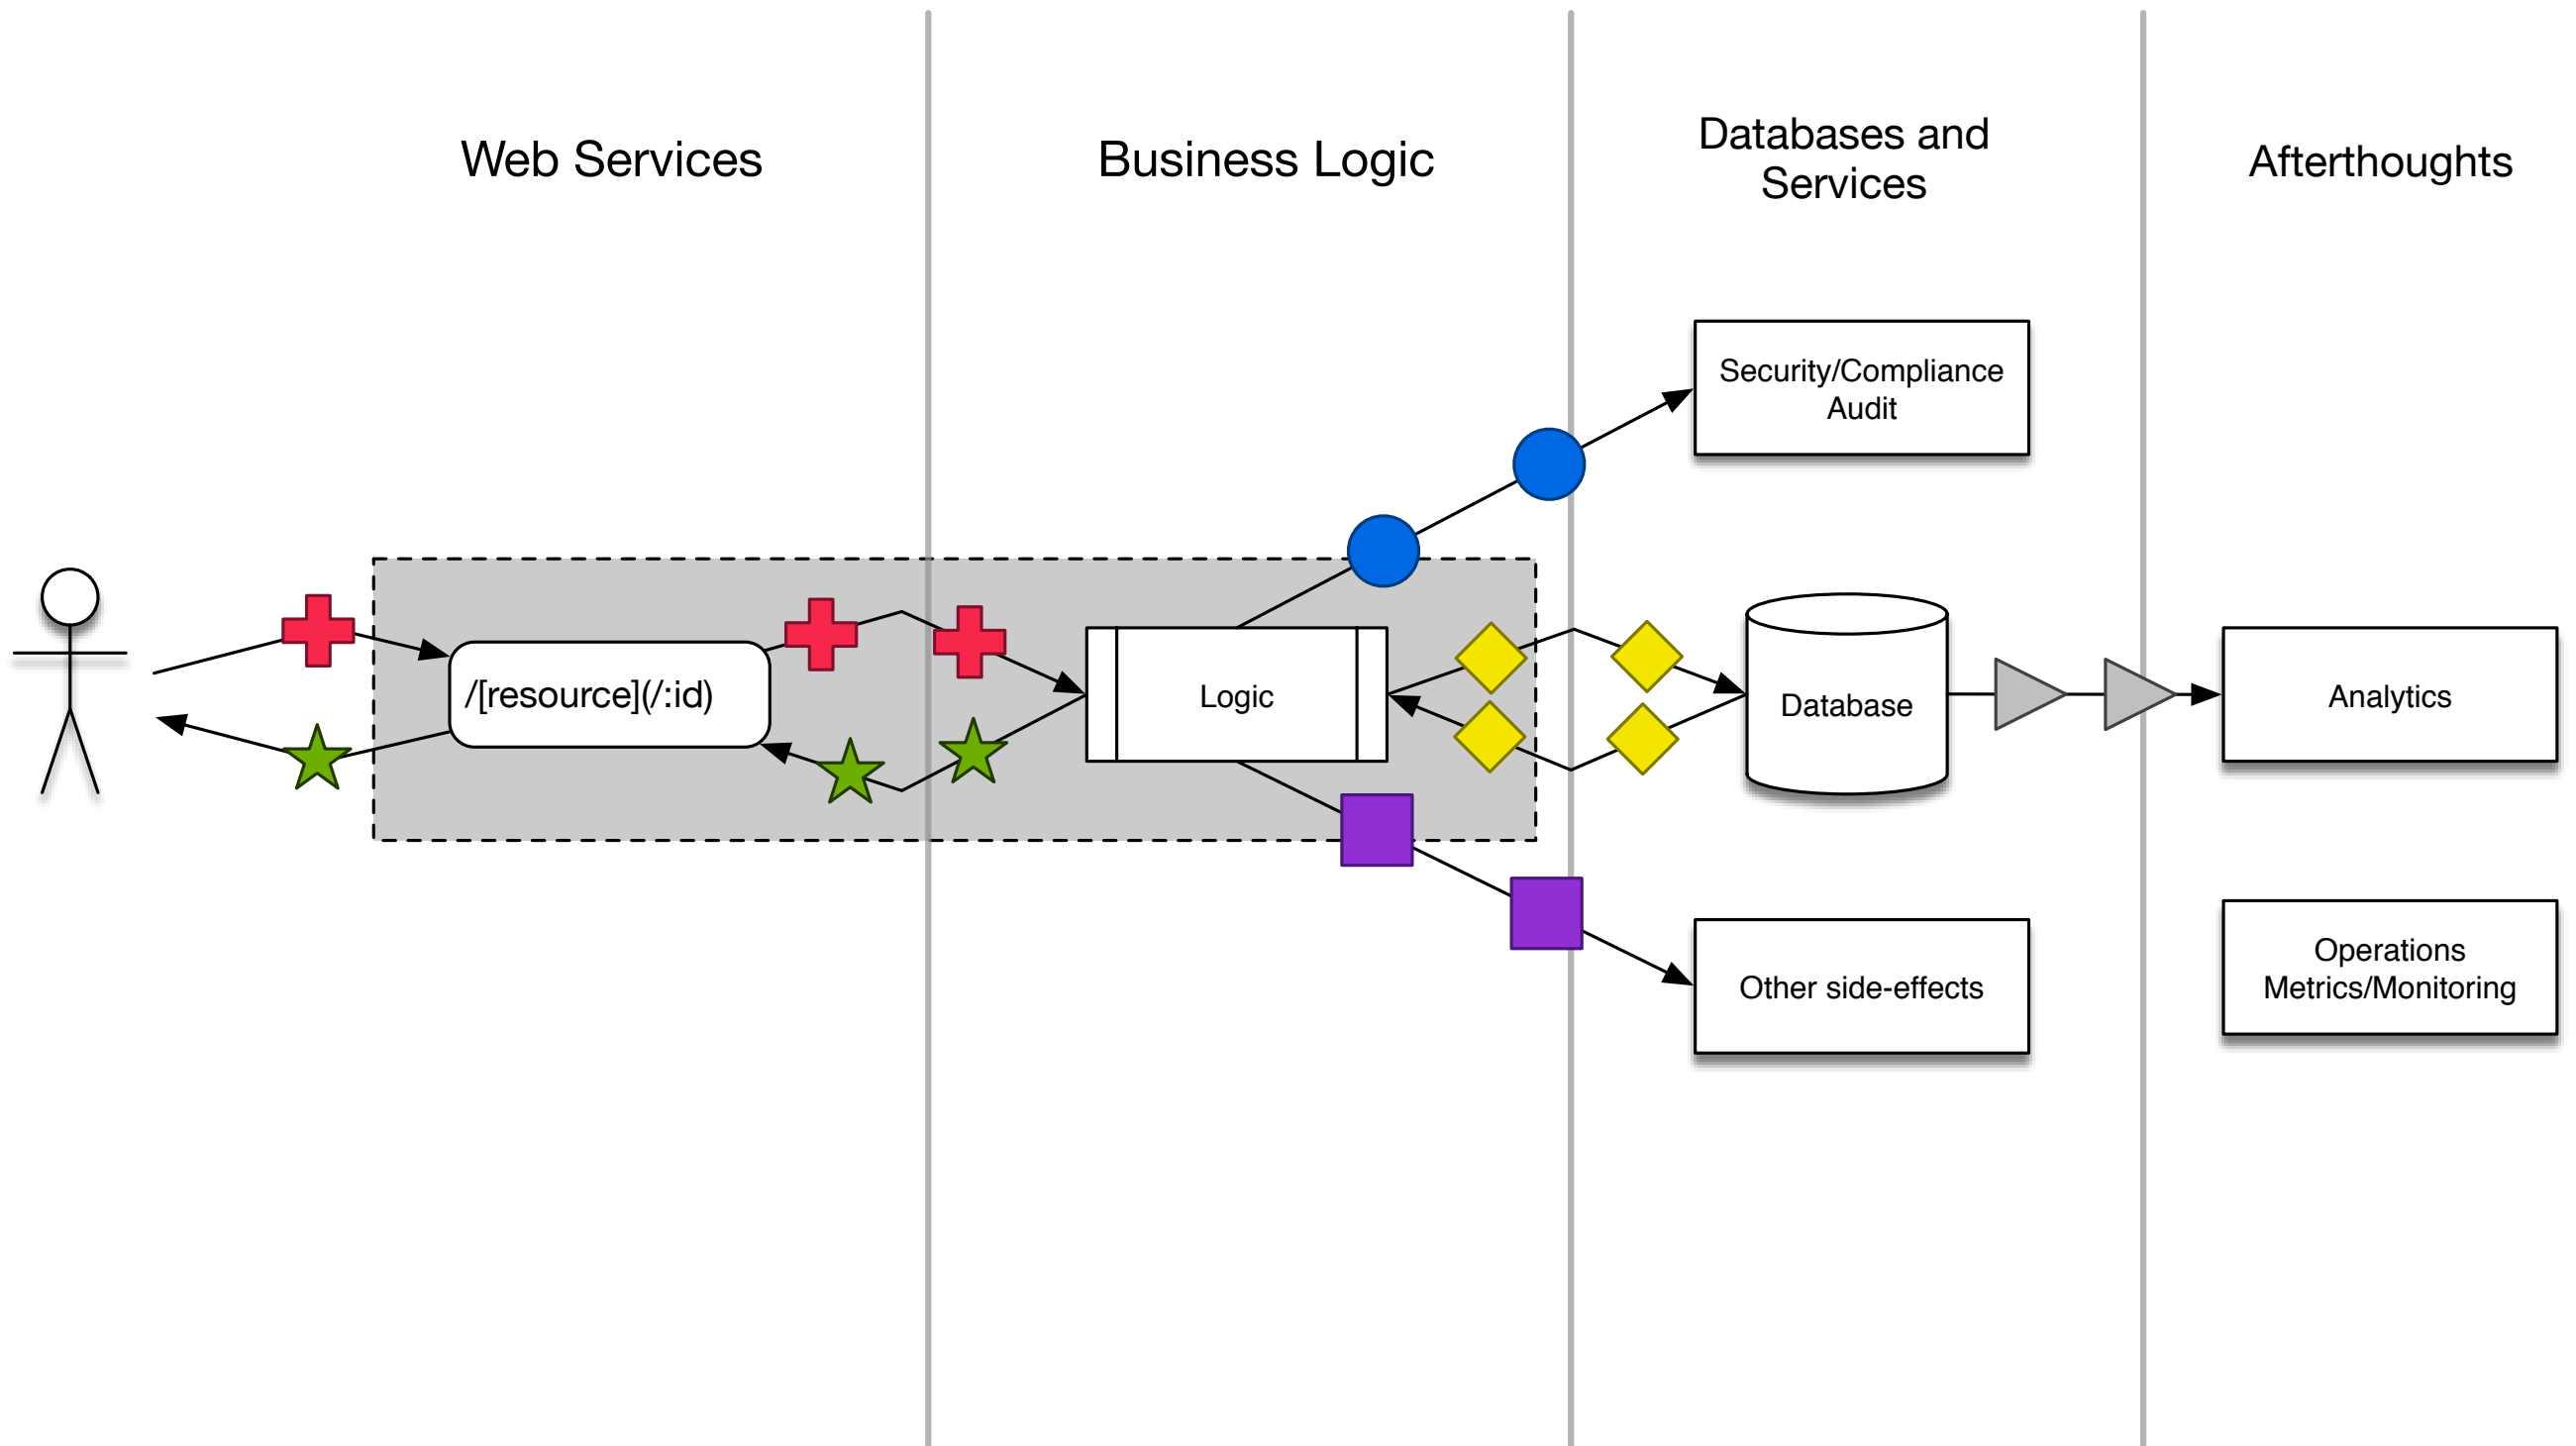


Business Services
are ***not*** Databases

“We shape our tools and thereafter our tools
shape us.”

–John M. Culkin

Database Leaking




Conway's Law

Yup, totally a thing

“organizations which design systems ... are
constrained to produce designs which are
copies of the communication structures of these
organizations”

–Melvin Conway

Cross-cutting Concerns?



We can do better!

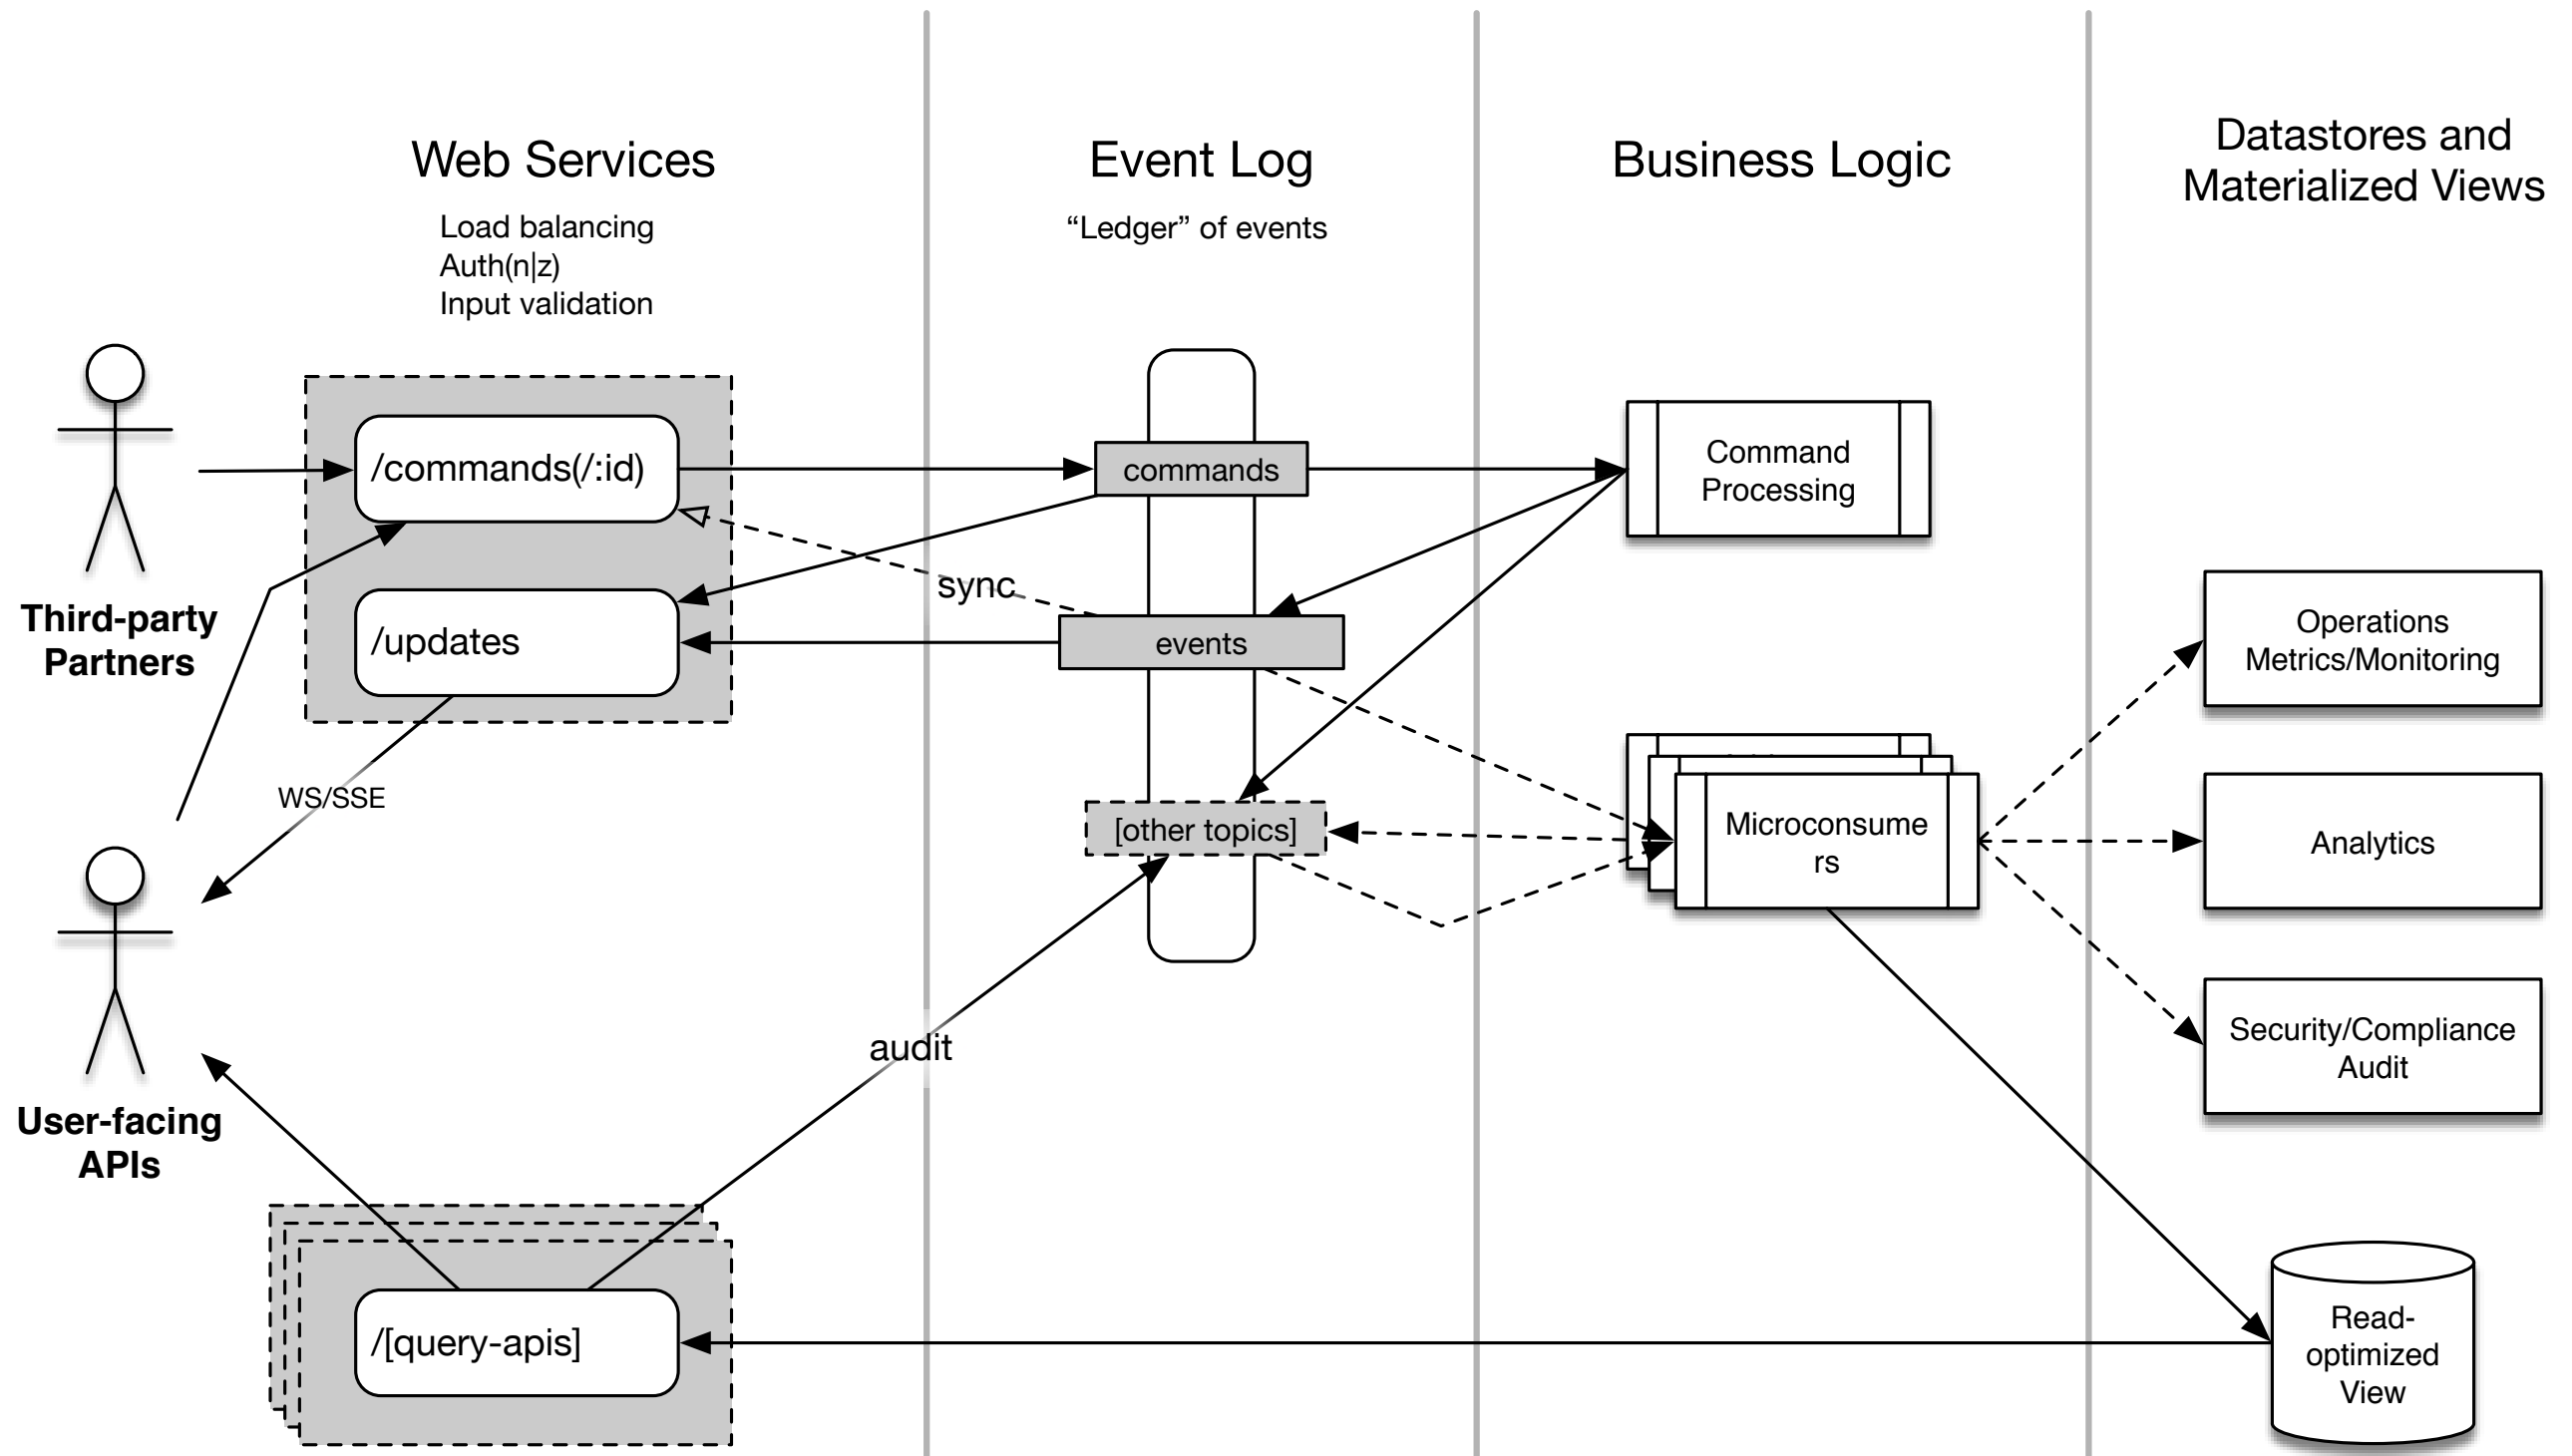
Commander

- A better architecture for APIs and services using REST + Immutable Event Log + Reactive Event Stream Processing
- The write-handling component of that architecture, my implementation is in Clojure

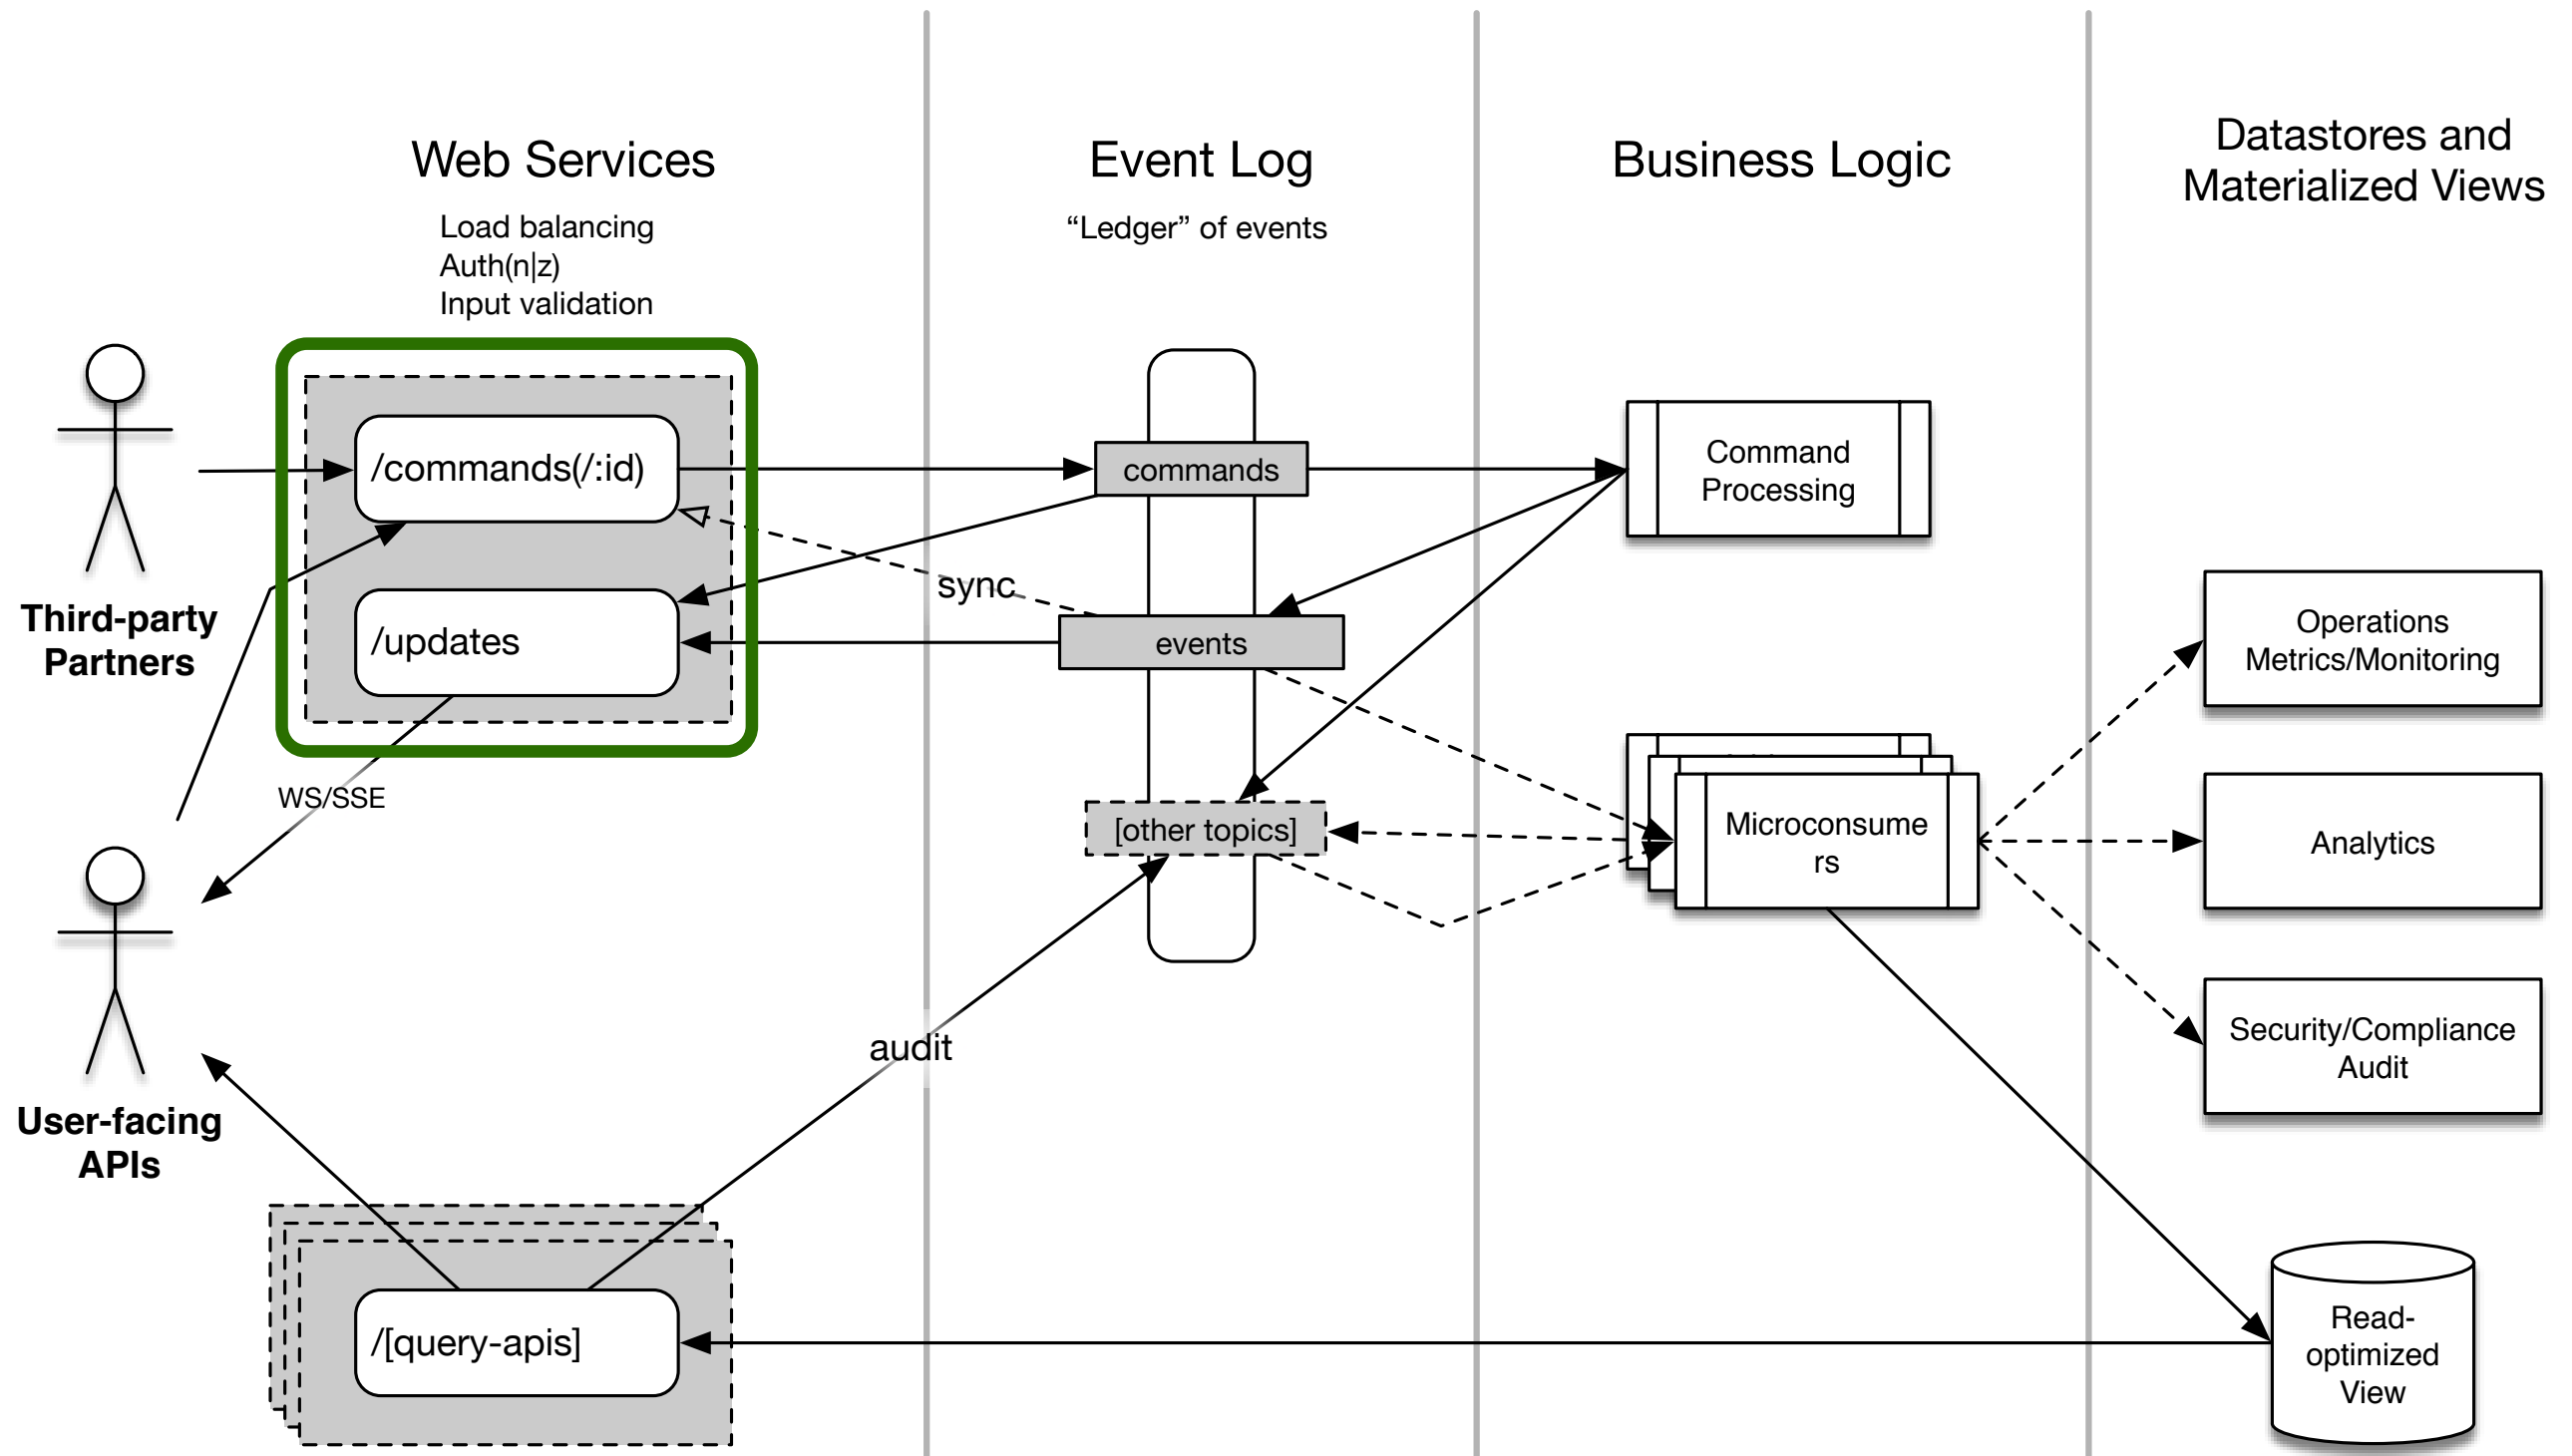
Commander Architecture

- Several categories of microservices with structured interactions among them
- REST + CQRS + Event Sourcing + Reactive Event Stream Processing

Commander Architecture

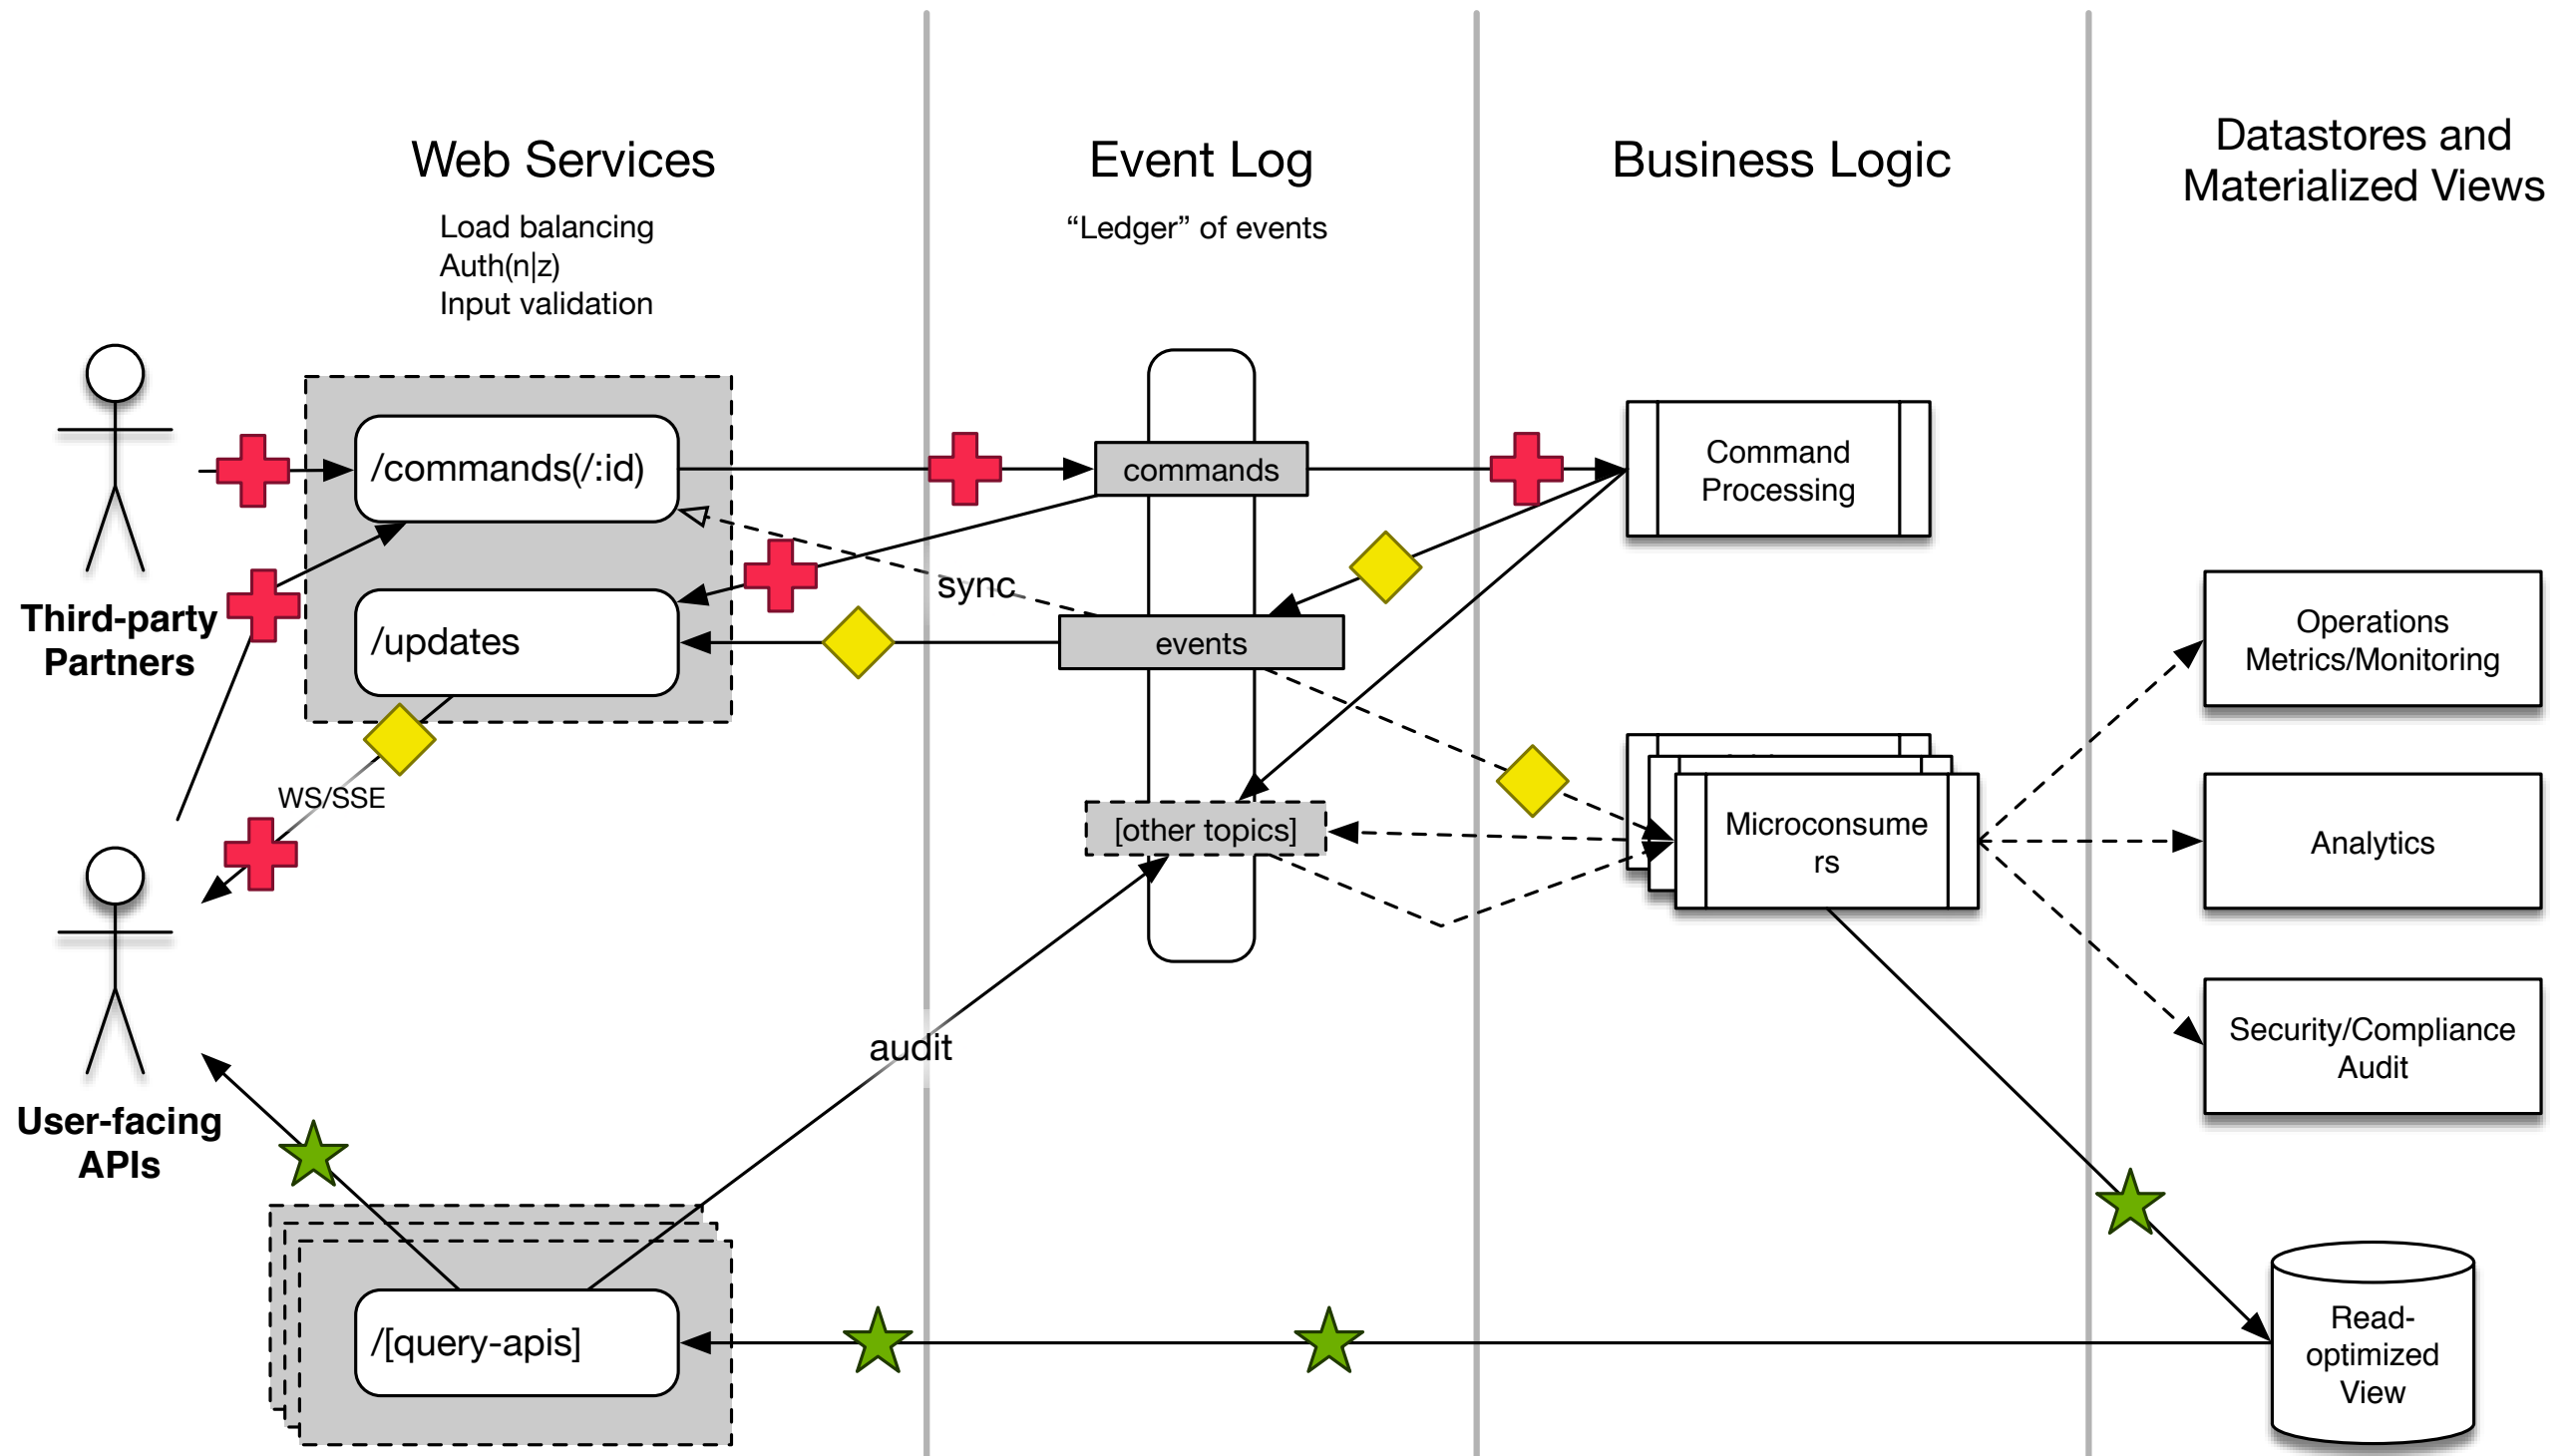


Commander Component




Embrace Immutability

Immutable Data Log




Express actions in
domain language
(not in database language)

A Command



```
{ "id": "33bb75db-6e13-48ee-8a54-b3976d3d065b",  
  "action": "transfer-money",  
  "data": { "from_account": "12345",  
            "to_account": "54321",  
            "amount": 10000 }  
  "timestamp": "2016-05-20T14:33:28.902-00:00" }
```

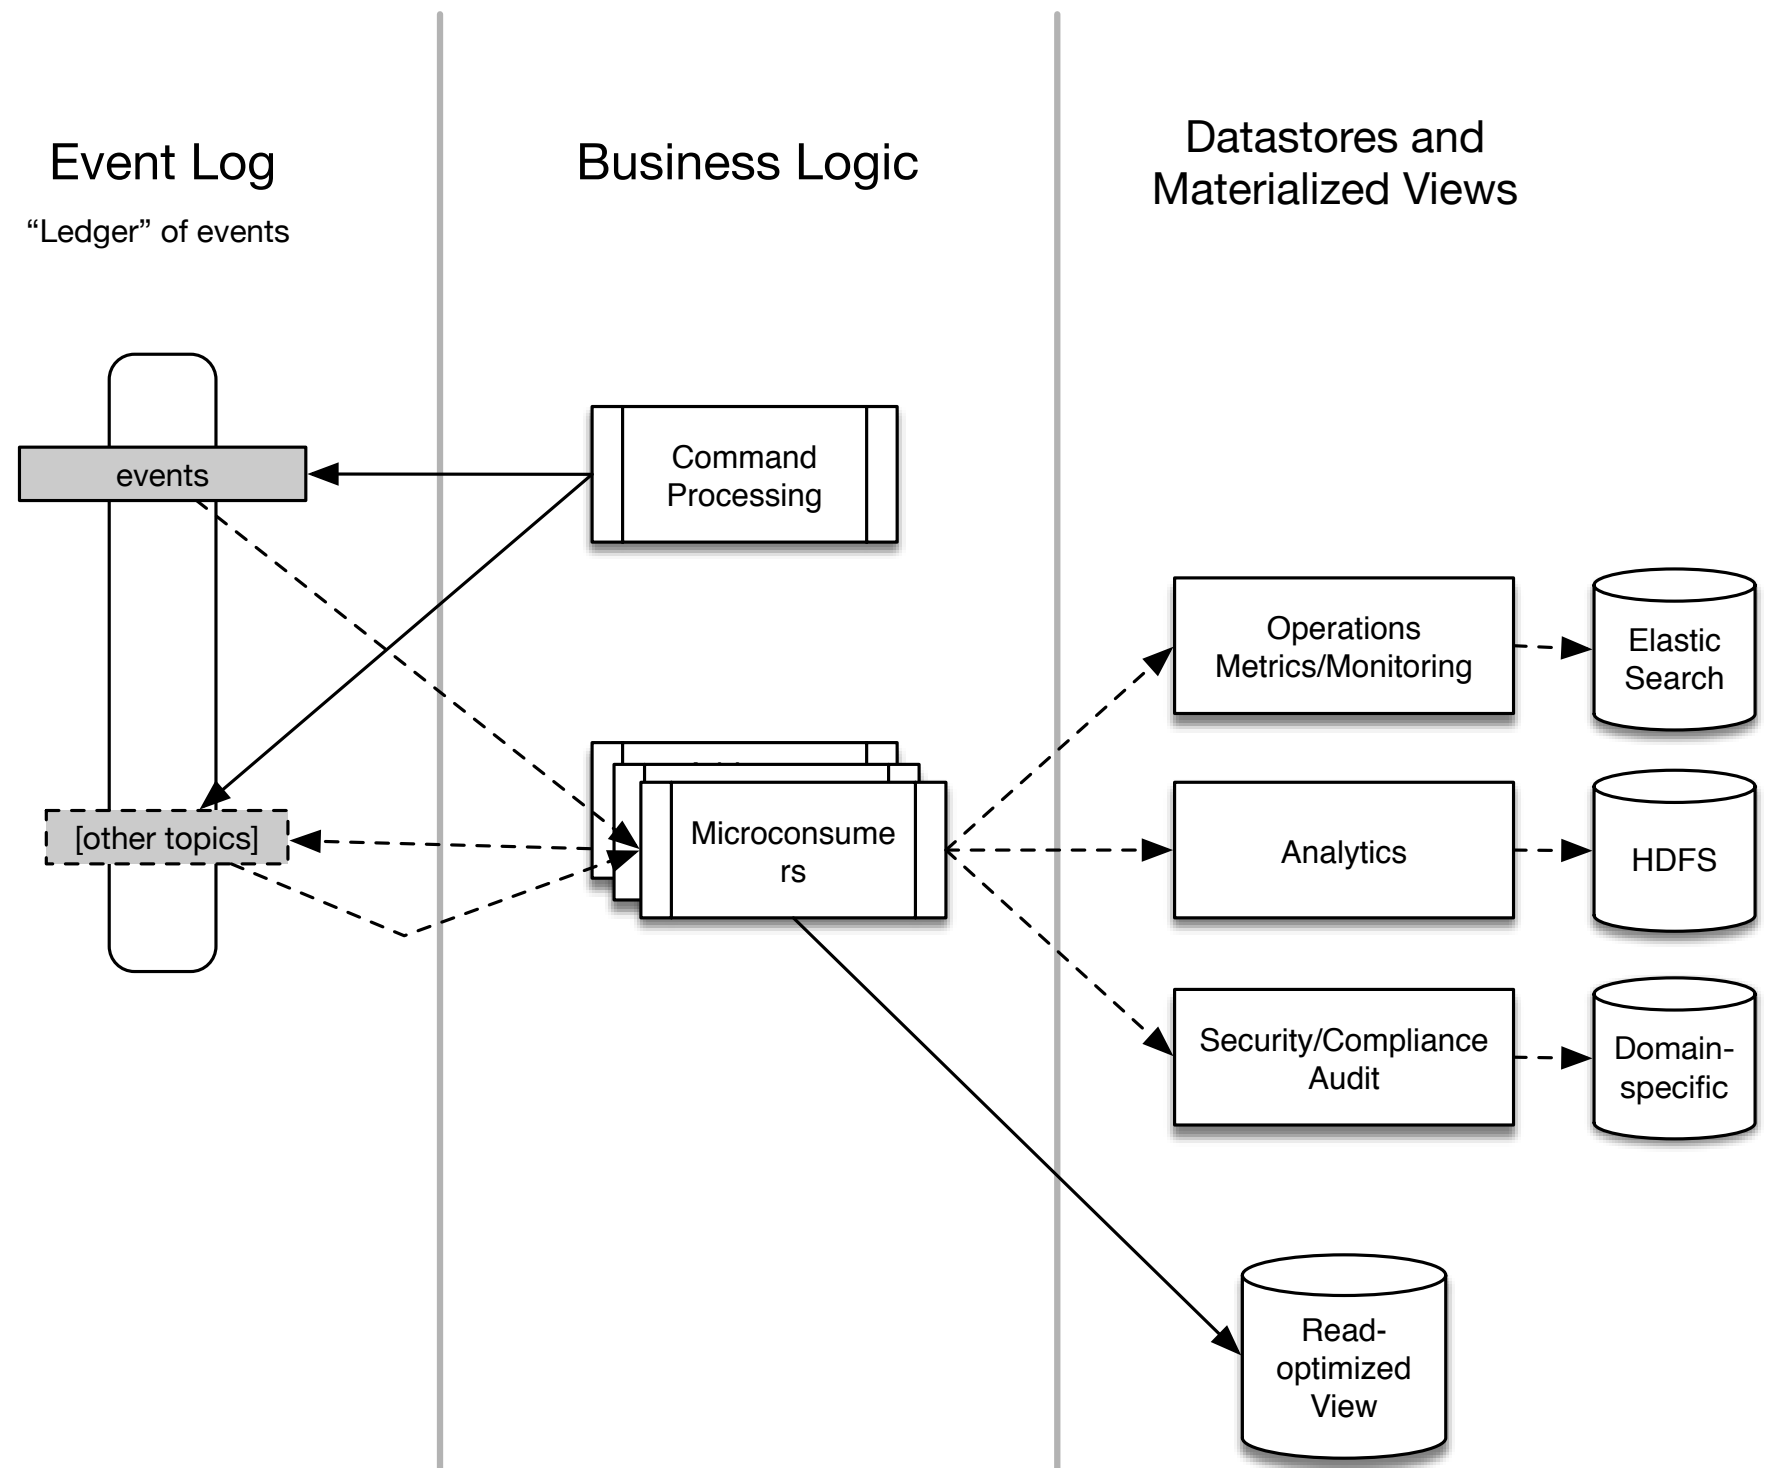
An Event



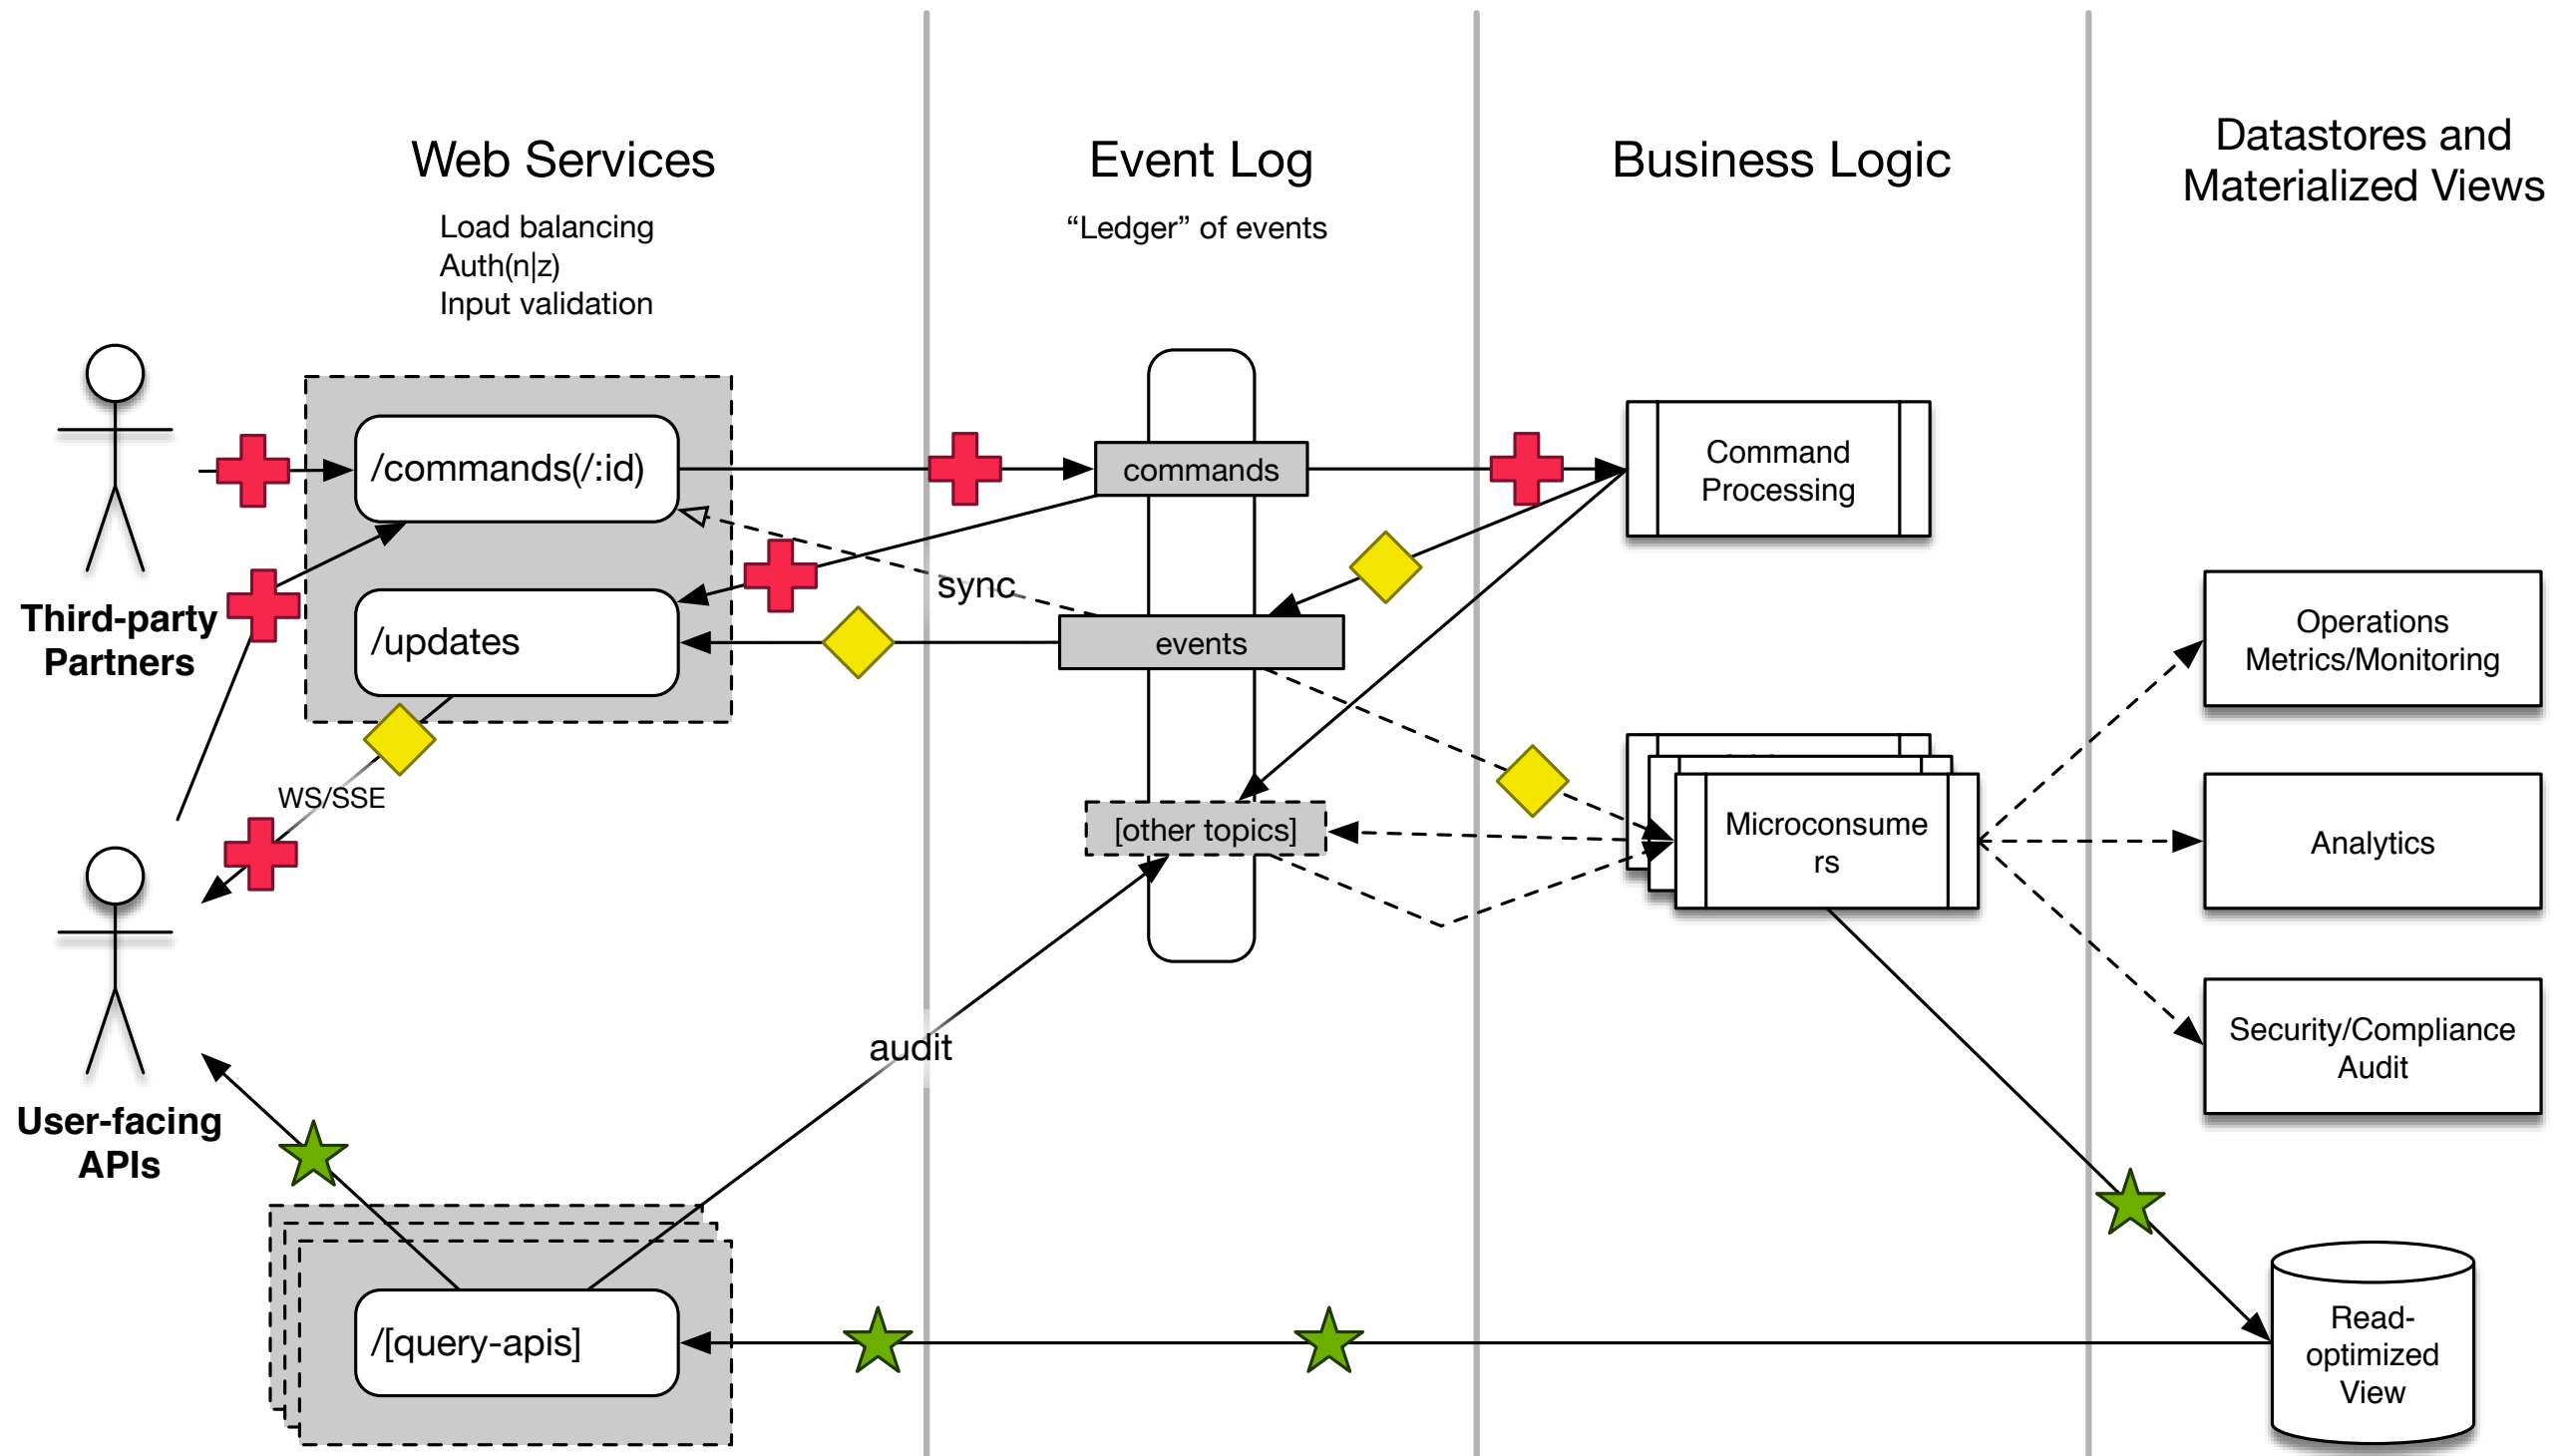
```
{ "id": "d435ed18-4ff7-4cae-a21b-3adb7b06fe58",  
  "parent": "33bb75db-6e13-48ee-8a54-b3976d3d065b",  
  "action": "money-transferred",  
  "data": { "id": "a6b903f6-0b9c-4c5b-95fa-afd4cc3bf938",  
            "from_account": "12345",  
            "to_account": "54321",  
            "amount": 10000 },  
  "timestamp": "2016-05-20T14:33:28.904-00:00" }
```

Separate Action
from Perception

1 Log => n Data Views

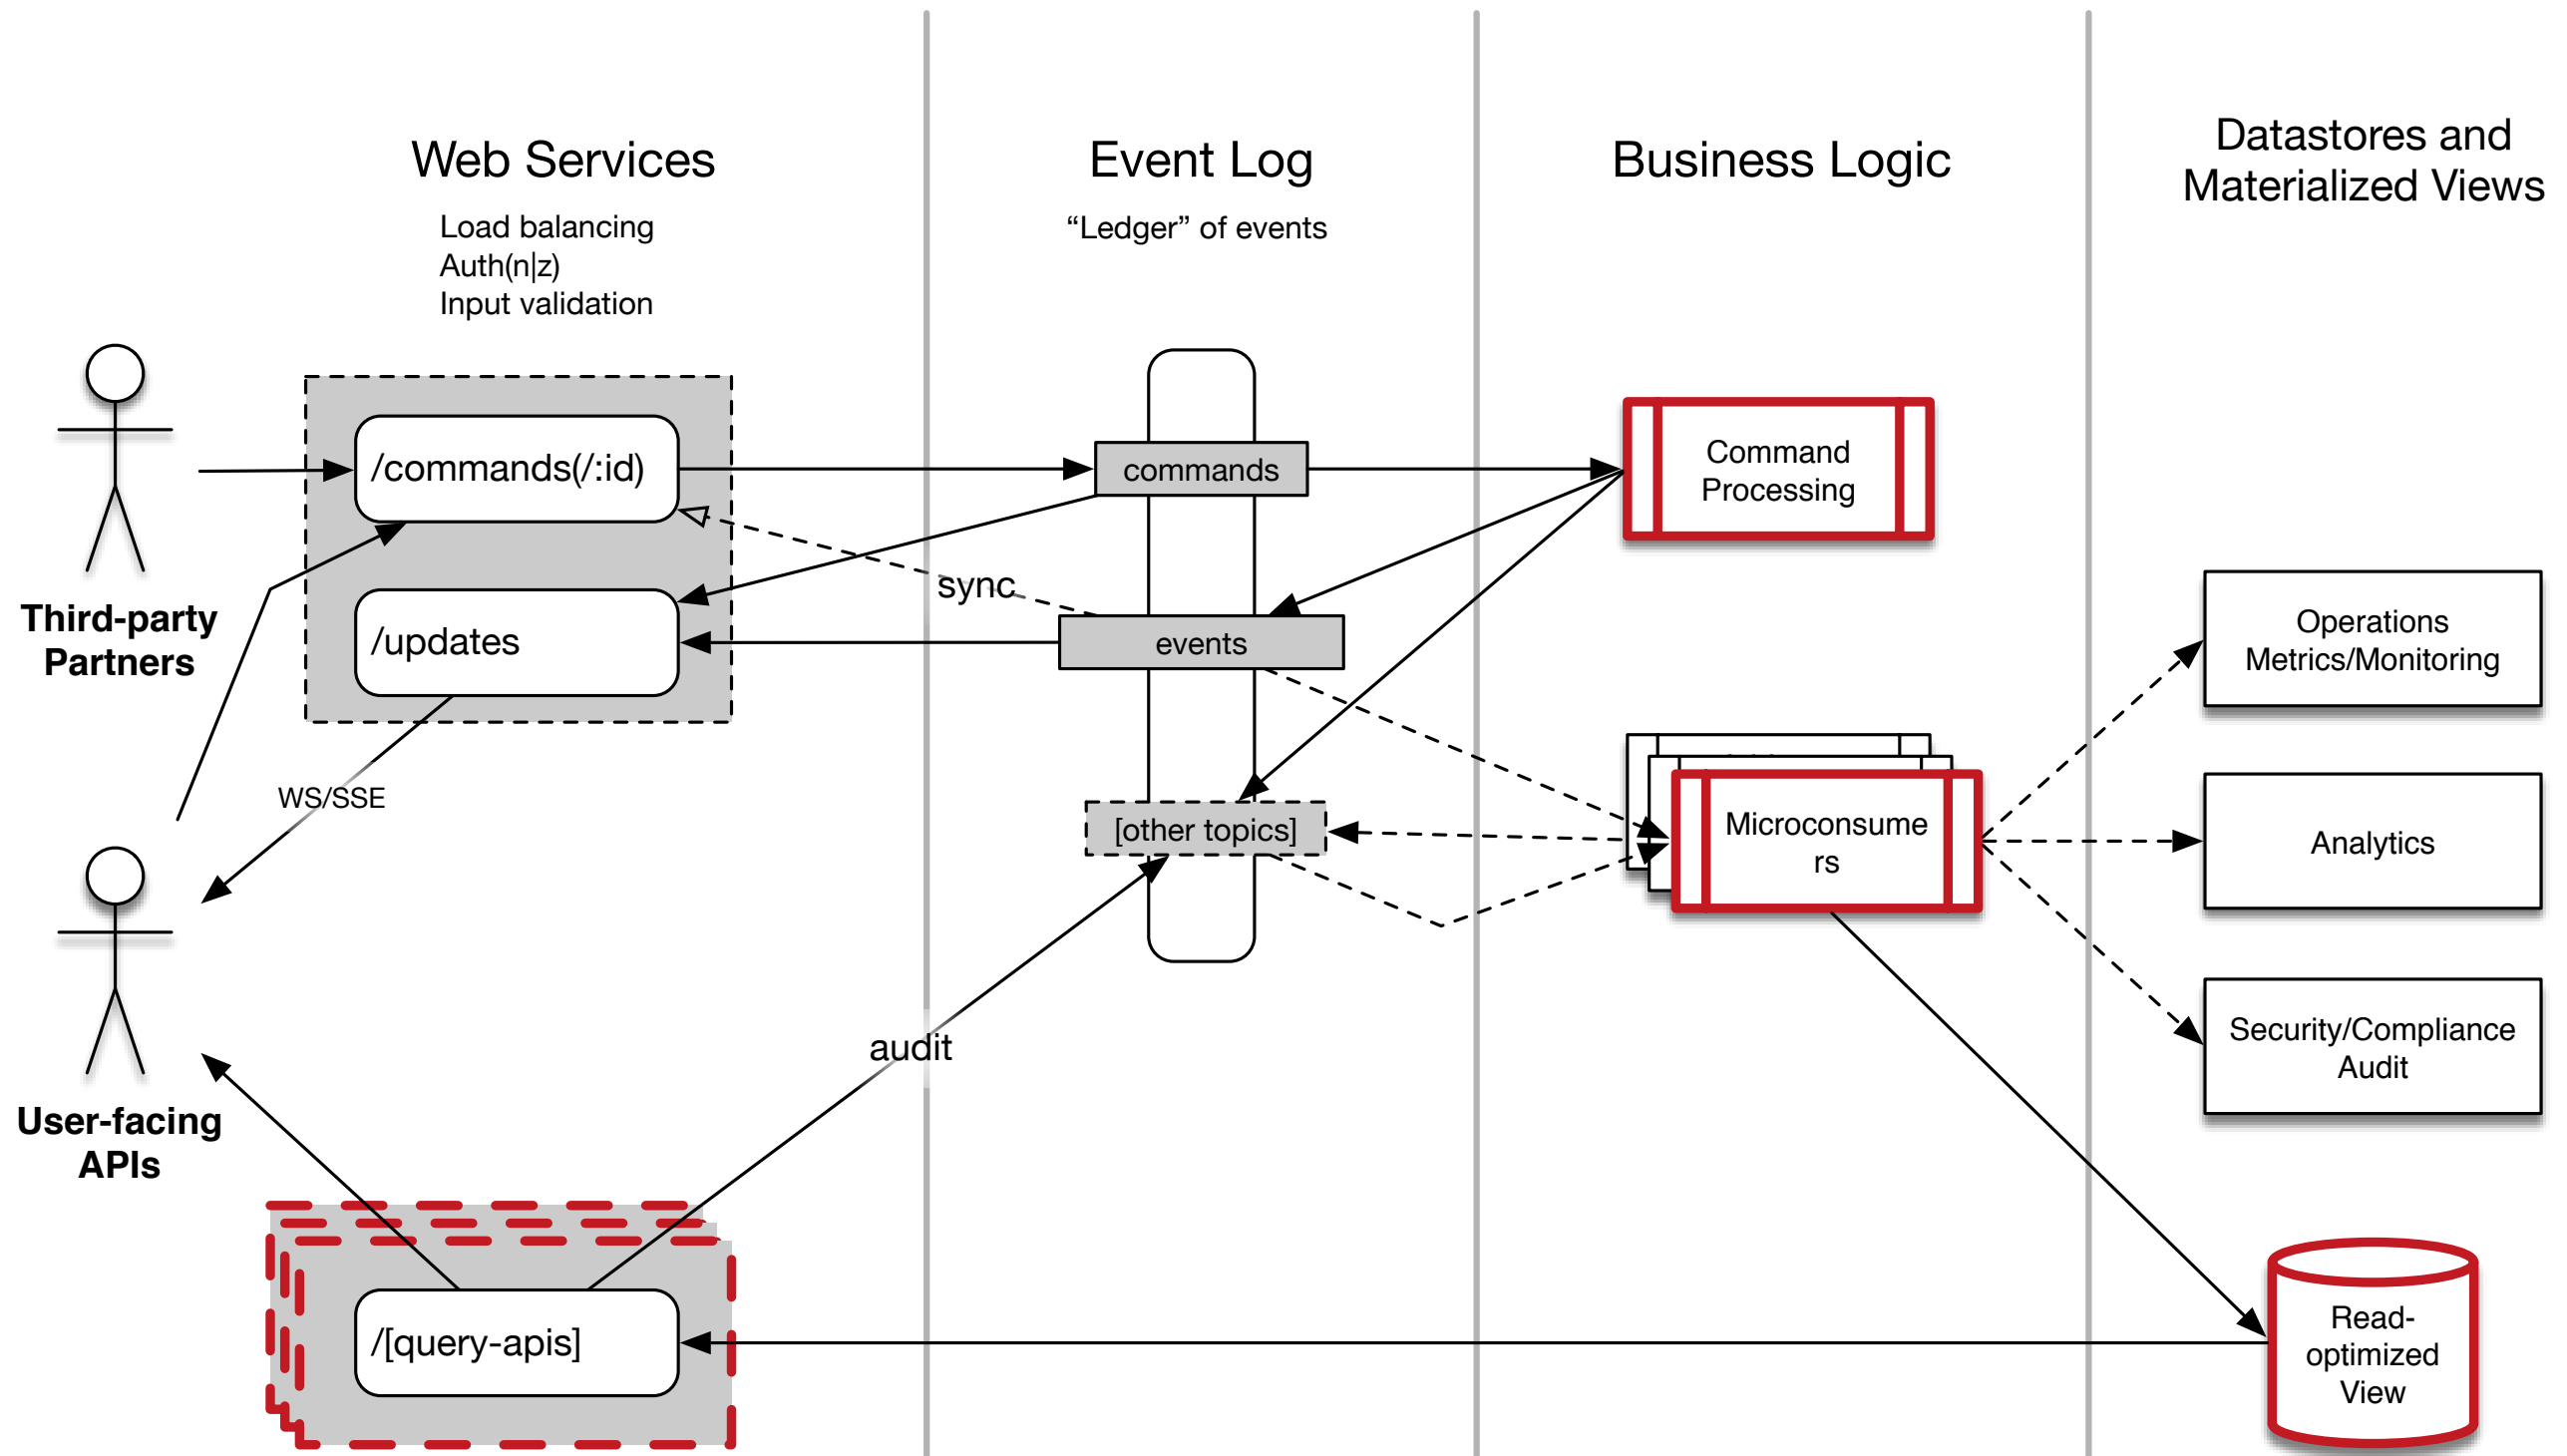


Commander Architecture

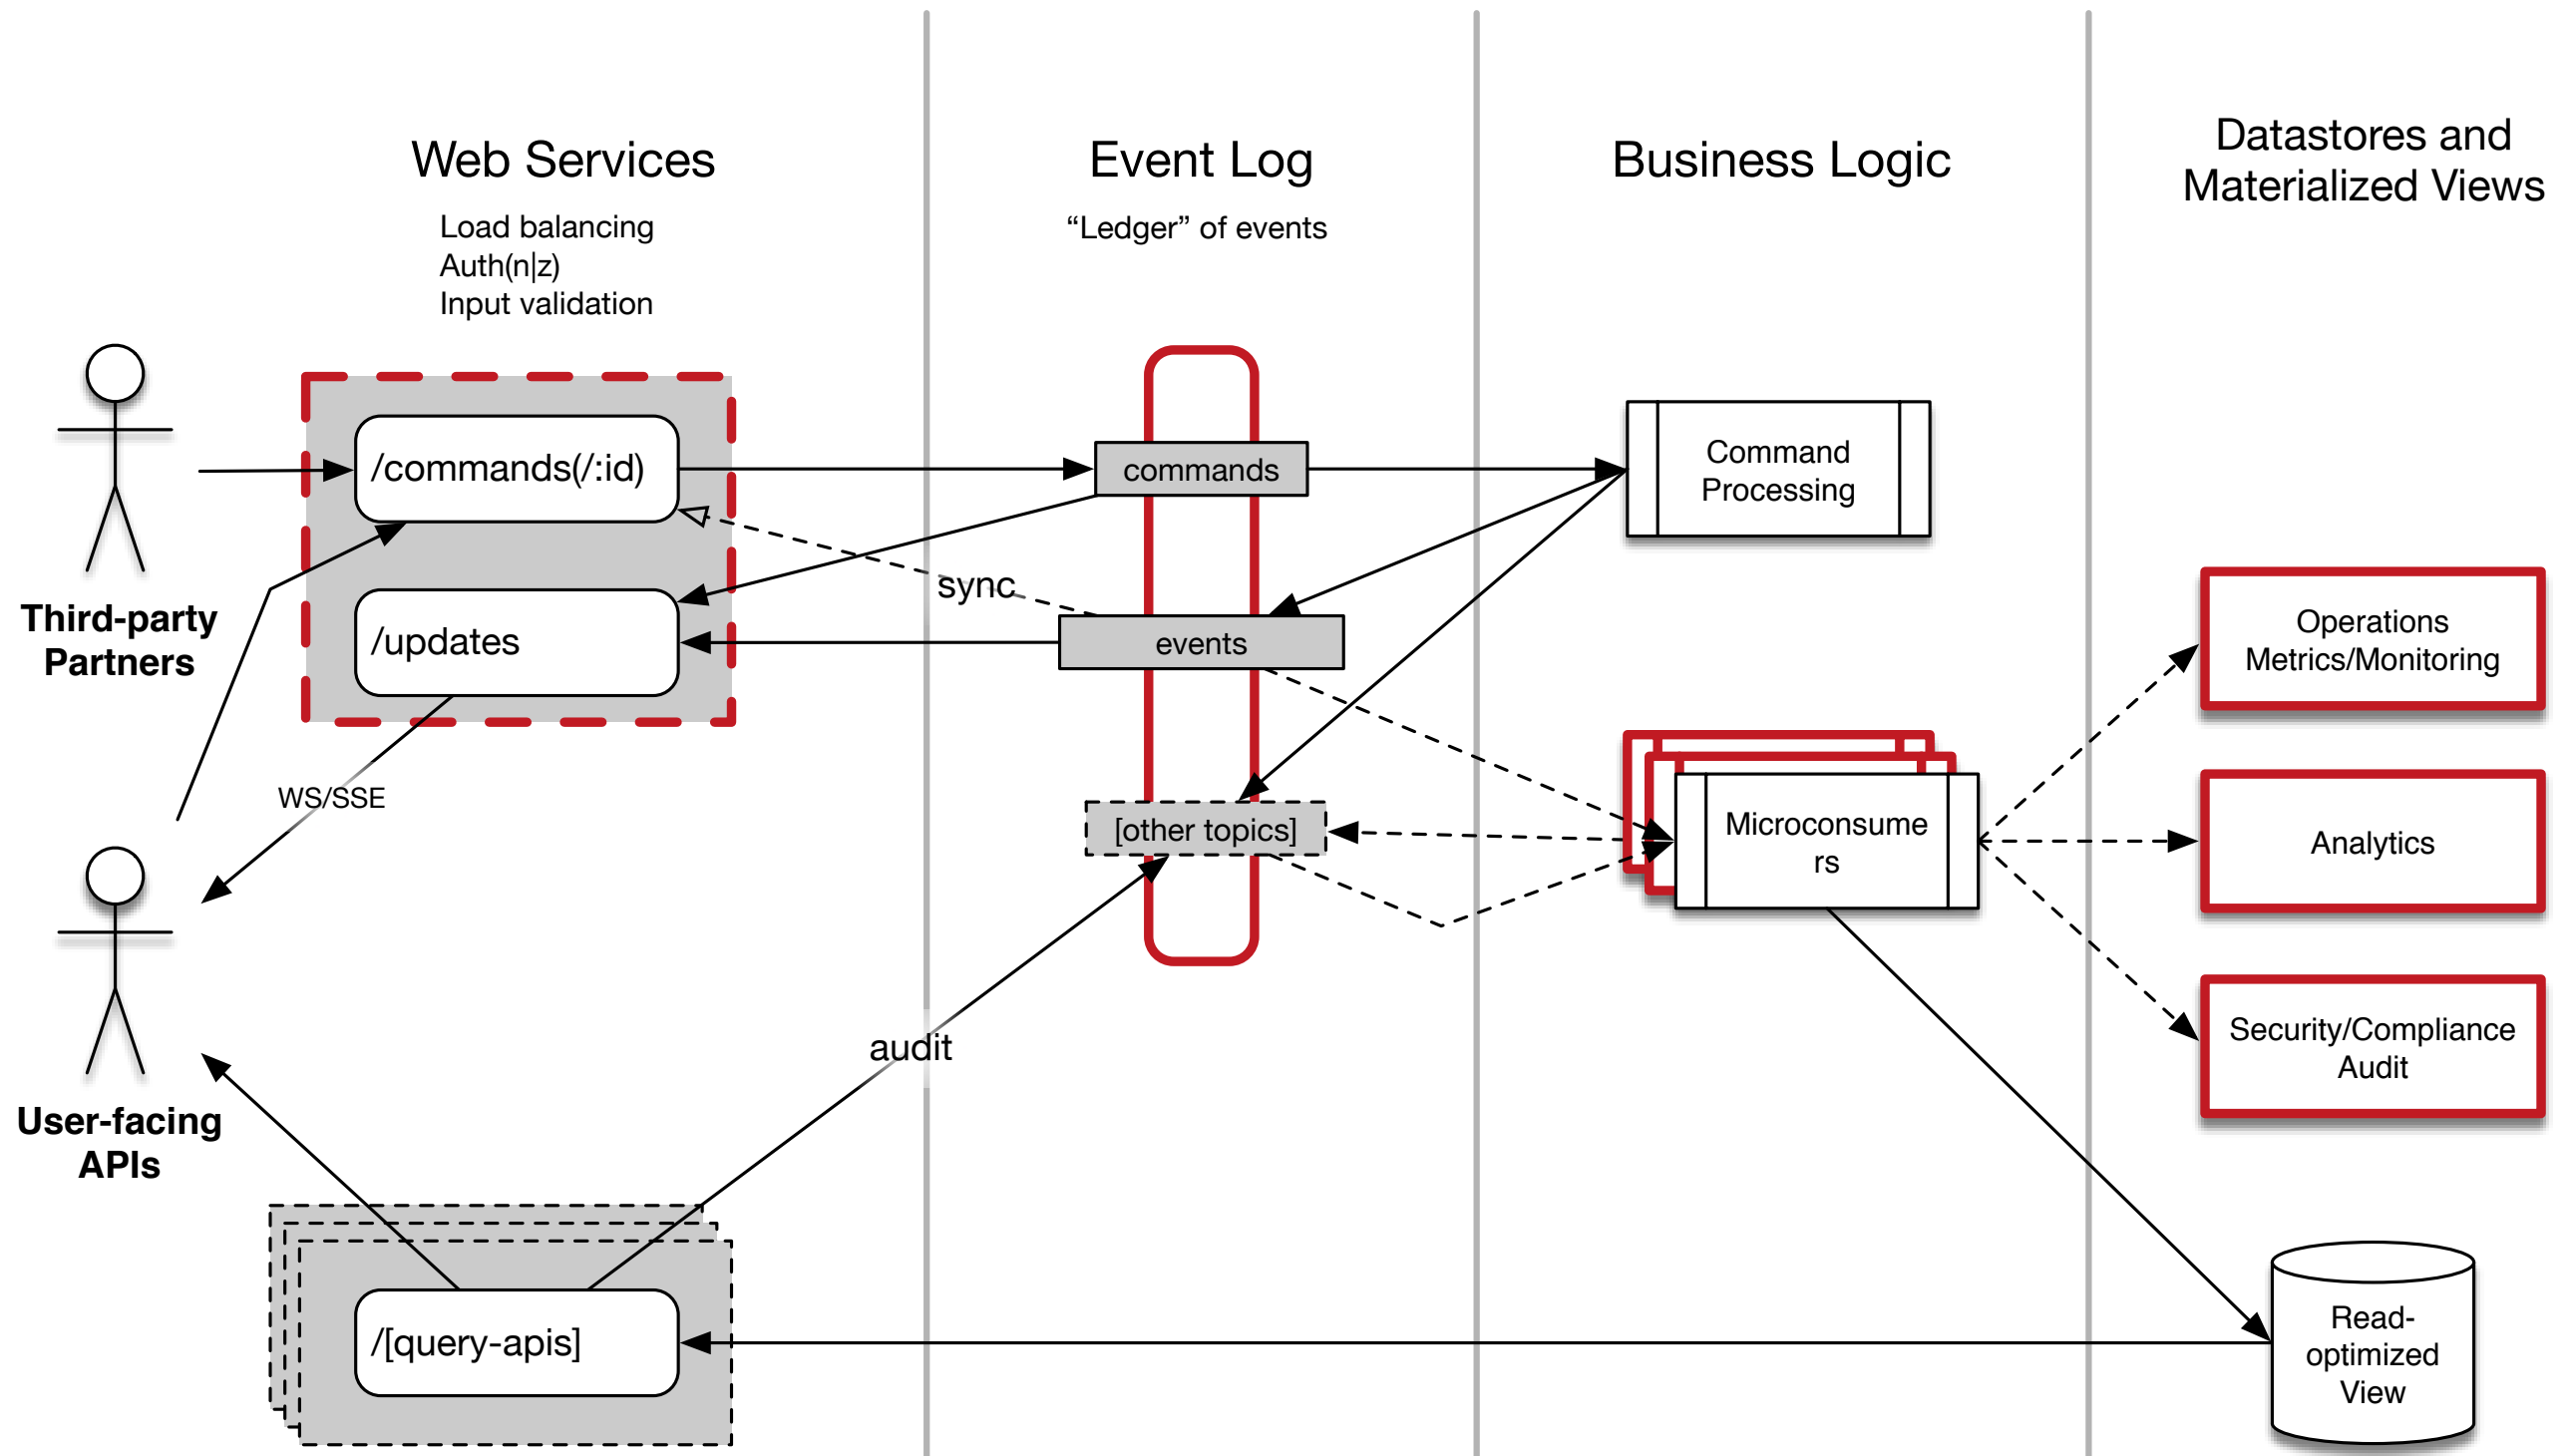


Exploit Conway's Law

Primary Team Provides



Enterprise Provides



Why Commander Component?

- Single writer to commands topic
- Ensuring schema conformance
- Indexing all Commands and Events for reads and server-push
- Provides optional illusion of synchrony to clients

How to implement
Reactive Services?

Kafka Streams!

What is Kafka?

- Apache Kafka is publish-subscribe messaging rethought as a distributed commit log
- But it's not really about messaging, that's just the interface
- Logs > Messages for my domain
- **It provides distributed, immutable logs!**


What is Kafka Streams?

- A Java library for building streaming applications on top of Kafka, lives in your application
- Low-level API for building topologies of processors, streams, and tables
- High-level DSL for common patterns like filter, map, aggregations, joins, stateful and stateless processing
- Nice operational characteristics (low latency, elastic, fault-tolerant)

How to use Kafka Streams within Commander

- Implement Command Processor
- Implement Event consumers and producers
- Provide local state management as backend for APIs

Commander Architecture



```
KStreamBuilder builder = new KStreamBuilder();

KStream<UUID, Map> commands = builder.stream(commandsTopic);
KStream<UUID, Map> customerEvents = commands
    .filter((id, command) -> command.get("action")
        .equals("create-customer"))
    .map((id, command) -> {
        Map userEvent = new HashMap(command);
        userEvent.put("action", "customer-created");
        userEvent.put("parent", id);
        Map userValue = (Map) userEvent.get("data");
        userValue.put("id", UUID.randomUUID());
        return new KeyValue<>(UUID.randomUUID(), userEvent);
    }).through(eventsTopic);

KStream<UUID, Map> customers = customerEvents
    .map((id, event) -> {
        Map customer = (Map) event.get("data");
        UUID customerId = (UUID) customer.get("id");
        return new KeyValue<UUID, Map>(customerId, customer);
    });

customers.through(customersTopic);

StateStoreSupplier store = Stores.create("Customers")
    .persistent()
    .build();
builder.addStateStore(store);

customers.process(customerStore, "Customers");

this.kafkaStreams = new KafkaStreams(builder, kafkaStreamsConfig);
this.kafkaStreams.start();
```


Demo

Summary

- Capture customer intent and business events as immutable data ***in domain language***
- From these action streams, services implement their own functionality in this common lingua franca
 - ***building many independent data views***
 - **reactively**
 - **without temporal or organizational coordination**

Giant Shoulders

- **Immutability:** Rich Hickey, Stu Halloway
- **CQRS:** Udi Dahan, Martin Fowler, Chris Richardson
- **Kafka Event Stream Reactors:** Neha Narkhede, Jay Kreps, Martin Kleppmann
- **Organization and Management:** Mel Conway, Eliyahu Goldratt, Gene Kim, Michael Nygard

References

- **<https://tinyurl.com/capital-one-cmdr>**
- <http://www.datomic.com/>
- <http://blog.cognitect.com/?tag=NewNormal+Series>
- <http://www.confluent.io/blog/event-sourcing-cqrs-stream-processing-apache-kafka-whats-connection/>
- <https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying>
- <https://www.infoq.com/presentations/Value-Values>